

CS 361: Introduction to Threads

Mark Boady

Department of Computer Science
Drexel University

September 28, 2022



Threads

- Most computers can run multiple programs at the same time
- Threads allow one program to do multiple things at the same time
- On a single processor: OS switches between tasks
- Multiple Cores or Processors: True concurrent execution

Threads

- Multiple things happening at once
- Threads can share memory
- Sometimes threads work independently
- Sometimes threads need to wait for other threads
- Sometimes threads need to access the same data at the same time

- Concurrent Tasks
 - 1 thread for GUI
 - 1 thread for physics simulator
 - 1 thread for database updates
- Parallel Processing
 - Break one problem into many smaller problems
 - Smaller problems must be solved independently
 - Hadoop for big data


- If two threads never communicate or share data, they truly run in parallel
- In most real programs, we need to share data or wait for results
- This means sometimes one thread is waiting on another
- The key is to avoid waiting unless required
- **Race Conditions:** something bad happens because two threads change memory at the same time
- **Deadlock:** something bad happens because two threads wait on each other endlessly

What is a thread?

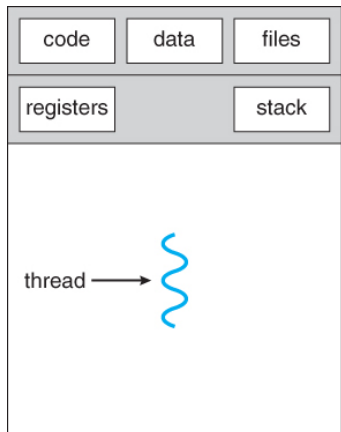
- Threads are partially independent and partially shared
- Shared:
 - All the program code loaded into memory
 - Optionally: any variable/object on the heap
 - We need to decide what variables/objects to share
- Independent:
 - Program Counter
 - Registers
 - Local Variables
 - Call Stack
- A thread should be **faster** than switching between two distinct programs due to the shared resources

Multiple Threads


- One program and span many threads
- Threads can begin and end at an point
- <https://www3.ntu.edu.sg/home/ehchua/programming/java/images/Multithread.gif>



Shared Memory



single-threaded process



multithreaded process

https://www.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/4_Threads.html

- C++ added threads to the standard in C++11
- C++ threads further improved in C++17
- We will use the features of C++17
- Even more concurrent features in C++20

- Part of the Standard Library
- Use `#include <thread>`
- Create a new thread `std::thread t(someFunction)`
- Wait for the thread to finish with `t.join()`
- Compile with special flags
`g++ -pthread --std=c++17 hello.cpp`
- *Note:* Not all OS need `-pthread`

- Span a new thread to print a message
- Wait till that thread finishes to exit the program

code/hello.cpp

```
15 //Include I/O stream for printing
16 #include <iostream>
17 //Include thread library
18 #include <thread>
19
20 /**
21  * This function prints a welcome message.
22  */
23 void welcome()
24 {
25     std::cout << "Welcome to CS361\n";
26 }
```

code/hello.cpp

```
33 int main(int argc, char** argv)
34 {
35     std::thread myThread(welcome);
36     //Wait for the thread to finish
37     myThread.join();
38     //Exit Program
39     return 0;
40 }
```

Creating a Thread

- We create a new thread with a constructor
- Syntax:
`std::thread threadName(functionToRun)`
- We pass in the name of the function we want the thread to execute
- We can also pass parameters to the function
- Syntax:
`std::thread threadName(func, in1, in2, in3)`

- There are multiple ways to share data between threads.
- Values are copied into the new thread by default
- Some Options to share data
 - Global Variables
 - Use Pointers
 - Use Explicit References

Example: Sleep Sort

- Sleep Sort is a classic (silly) algorithm to show threads
- For each item, make a thread that sleeps that many seconds
- Each thread prints out its number when it wakes up.

Sleep Sort Algorithm

```
function SLEEP_THREAD(value)  
    Sleep for value seconds  
    Print Value  
return  
end function
```

Sleep Sort Algorithm

```
function SLEEPSORT(Array A, int size)
  for i=0; i < size; i++ do
    sleepThread(A[i])
  end for
  Wait for all threads to end
end function
```

Waiting for a Thread

- How do we wait for a thread?
- `t.join()` waits until this thread has exited
- What if we don't want to wait?
- `t.detach()` frees the thread to run, we won't wait for it
- A detached thread is still tied to main, it exits with the program

Initial Setup

- We want to generate random arrays
- The size of the array comes from the command line
- We want to print the array so we can see it worked
- None of this need threads.

code/sleep_01.cpp

```
18  /**
19   Generate an array of random numbers between 0
      and n*5
20   @param n is the size of the array to create
21   @return Pointer to the array created
22   */
23  int* randomArray(int n);
24  /**
25   Print out an array that is nicely formatted
26   @param A is a pointer to the array
27   @param size is the number of elements in the
      array
28  */
29  void printArray(int* A, int size);
```

code/sleep_01.cpp

```
31 /**  
32  Make an array with the given number of  
    elements. Sort using sleepy sort.  
33  @param argc is the number of command line  
    arguments  
34  @param argv is the text of the command line  
    arguments  
35  @return 1 on failure and 0 on success  
36  */
```

code/sleep_01.cpp

```
37 int main(int argc , char** argv)
38 {
39     //Check we got a command line argument
40     if(argc!=2)
41     {
42         std::cout <<
43             "Usage: sleepy [number elements]"
44             << std::endl;
45         return 1;
46     }
```

code/sleep_01.cpp

```
47 //Use the command line argument as the
    size
48 int size = std::atoi(argv[1]);
49 //Initiate a random number generator
50 std::srand(std::time(NULL));
51 //Make an array
52 int* myArray = randomArray(size);
53 //Print the array to see what it looks
    like
54 printArray(myArray, size);
55 //Return and exit program
56 return 0;
57 }
```


code/sleep_01.cpp

```
61 int* randomArray(int n){
62     int* A = new int[n];
63     for(int i=0; i < n; i++)
64     {
65         A[i] = std::rand()%(n*5);
66     }
67     return A;
68 }
```

code/sleep_01.cpp

```
70 void printArray(int* A, int size)
71 {
72     std::cout << "[";
73     for(int i=0; i < size; i++)
74     {
75         std::cout << A[i];
76         if(i+1!=size)
77         {
78             std::cout << ", ";
79         }
80     }
81     std::cout << "]" << std::endl;
82 }
```

Sleep Sort

- We need a function for the threads to run.
- Each thread will:
 - 1 Wait x seconds
 - 2 Print it's ID
 - 3 Print the value x
 - 4 End

We need new includes and function prototypes

code/sleep_02.cpp

```
24  /**
25   Thread that sleeps for our sort
26   @param value is the number to sort
27  */
28  void sleepyThread(int value);
29  /**
30   Sort an array using Sleep Sort
31   @param array is the array to sort
32   @param size is the number of elements in the
        array
33  */
34  void sleepSort(int* array, int size);
```

Implementation of the Sleep Thread (Part 1)

code/sleep_02.cpp

```
105 void sleepyThread(int value){
106     //What is the Unique ID of this thread
107     std::thread::id myID =
108         std::this_thread::get_id();
109     //Determine how long to wait
110     std::chrono::seconds waitTime
111         = std::chrono::seconds(value);
```

Implementation of the Sleep Thread (Part 2)

code/sleep_02.cpp

```
112 // Fall Asleep
113 std::this_thread::sleep_for(waitTime);
114 // Print Value
115 std::cout << "Thread "
116           << myID << " says "
117           << value << std::endl;
118 // Exit
119 return;
120 }
```

Sleep Sort Implementation

- We need to join all threads, otherwise the program will exit before we are done.
- We can't join in the loop we construct the threads
 - The loop won't continue until the first thread ended
 - No real parallel processing!
- Store threads in an array and join after all are started

Implementation of the Sleep Sort (Part 1)

code/sleep_02.cpp

```
123 void sleepSort(int* A, int size)
124 {
125     //Make an Array of threads to check
        against
126     std::thread* myThreads = new std::thread[
        size];
127     //Generate the threads
128     for(int i=0; i < size; i++)
129     {
130         myThreads[i] = std::thread(
            sleepyThread, A[i]);
131     }
```


Implementation of the Sleep Sort (Part 2)

code/sleep_02.cpp

```
132 //Wait for all the threads to end
133 for(int i=0; i < size; i++)
134 {
135     myThreads[i].join();
136 }
137 //Exit Function
138 return;
139 }
```

Example Execution

```
1 [22, 28, 46, 38, 39, 10, 13, 39, 47, 47]
2 Thread 0x70000995c000 says 10
3 Thread 0x7000099df000 says 13
4 Thread 0x7000096cd000 says 22
5 Thread 0x700009750000 says 28
6 Thread 0x700009856000 says 38
7 Thread Thread 0x7000098d9000 says 0
   x700009a62000 says 39
8 39
9 Thread 0x7000097d3000 says 46
10 Thread Thread 0x700009ae5000 says 0
    x700009b68000 says 4747
```

Race Condition

- The print statements overlapped with each other!
- Both threads were trying to use cout at the same time
- **Race Condition:** when two or more threads try to access the same memory at the same time
- Bad things can happen!

- A **lock** is a data structure that ensures only one thread accesses something at a time
- We can have multiple **locks** to manage different resources
- New problems:
 - **Deadlock**: When a series of threads are all stuck waiting for each other to release a lock. It becomes impossible for any thread to move forward.
 - Slowdown is introduced any time a thread needs to wait its turn
- Too much code protected by locks forces everything to run in sequence

Locked Function

- We can lock an entire function
- Only one thread will be able to use this function at a time
- The lock is automatically released when the function returns
- Makes it easier to organize and manage locks
- We need shared lock for all the threads

Updated Include and global

code/sleep_03.cpp

```
19 //Needed for threading
20 #include <thread>
21 //Needed for sleep
22 #include <chrono>
23 //Needed for mutex lock
24 #include <mutex>
25 //Make our global lock
26 std::mutex coutLock;
```

New Print Function with Lock

code/sleep_03.cpp

```
154 void printLock(std::thread::id id, int num){
155     //Create a lock that lasts the
156     //life of this function
157     const std::lock_guard<std::mutex>
158         lock(coutLock);
159     //Print to Cout
160     std::cout << "Thread "
161         << id << " says "
162         << num << std::endl;
163 }
```

Example Execution

```
1  [8, 7, 8, 0, 8, 3, 4, 1, 2, 3]
2  Thread 0x700002e17000 says 0
3  Thread 0x700003023000 says 1
4  Thread 0x7000030a6000 says 2
5  Thread 0x700002f1d000 says 3
6  Thread 0x700003129000 says 3
7  Thread 0x700002fa0000 says 4
8  Thread 0x700002d11000 says 7
9  Thread 0x700002c8e000 says 8
10 Thread 0x700002d94000 says 8
11 Thread 0x700002e9a000 says 8
```


- The values are sorted in order.
- They are not sorted in the array!
- We can have the thread store in the array instead of printing
- We need to know what index to store to
- The index variable needs to be locked, we don't equal numbers in the same index

Array Index

- Each Thread will update a different index in the array
- The array does not need to be locked, no threads will overlap
- The index needs to be updates by each thread
- The index needs to be locked or two threads might look at the wrong index
- The index lock protects the whole array
- We lock inline to protect only the lines we need
- Counter is passed by reference

Revised Includes

code/sleepsort.cpp

```
12 //cout and endl
13 #include <iostream>
14 //We need srand and rand
15 #include <cstdlib>
16 //We can use the time to set the random
   generator
17 #include <ctime>
18
19 //Needed for threading
20 #include <thread>
21 //Needed for sleep
22 #include <chrono>
23 //Needed for mutex lock
24 #include <mutex>
```

Update to Sleep Sort Algorithm (Part 1)

code/sleepsort.cpp

```
141 void sleepSort(int* A, int size){
142     int index=0;
143     //Array of threads
144     std::thread* myThreads
145         = new std::thread[size];
146     //Generate the threads
147     for(int i=0; i < size; i++)
148     {
149         myThreads[i] = std::thread(
150             sleepyThread ,
151             A, //Array to change
152             A[i], //value to sleep on
153             std::ref(index)); //target
154     }
```

Update to Sleep Sort Algorithm (Part 2)

code/sleepsort.cpp

```
154 //Wait for all the threads to end
155 for(int i=0; i < size; i++)
156 {
157     myThreads[i].join();
158 }
159 //Exit Function
160 return;
161 }
```

Revised Sleep Thread (Part 1)

code/sleepsort.cpp

```
120 void sleepyThread(int * array, int value, int  
    &index){  
121     //Determine how long to wait  
122     std::chrono::seconds waitTime  
123         = std::chrono::seconds(value);  
124     //Fall Asleep  
125     std::this_thread::sleep_for(waitTime);  
126     //Update Array  
127     int myIndex;
```

Revised Sleep Thread (Part 2)

code/sleepsort.cpp

```
128 //Lock so only this thread can change
129 counterLock.lock();
130 //Update Index
131 myIndex = index++;
132 //Free Lock
133 counterLock.unlock();
134 //Update array
135 array[myIndex]=value;
136 //Exit
137 return;
138 }
```

Example Execution

1	[9, 7, 9, 8, 8, 5, 2, 8, 4, 4]
2	[2, 4, 4, 5, 7, 8, 8, 8, 9, 9]

- Threads can be created
- We can join or detach threads
- Values can be shared between threads
- Race Conditions happen when threads use the same memory at the same time
- Locks and protect memory
- Locks can cause problems (more on this in future weeks)